

LINEAR MODELS FOR CLASSIFICATION (PART 1)

Hsin-Min Lu

盧信銘

台大資管系

Topics

- Introduction
- Discriminant Function
- Probabilistic Discriminative Models
- Laplace Approximation
- Bayesian Logistic Regression



INTRODUCTION

Introduction

- Regression:
 - Map input vector x to one or more continuous target variables t
- Classification:
 - Map input vector x to one of K discrete classes ($K \geq 2$)
- Ordinal Regression:
 - Map input vector x to one of K discrete classes that have an ordering (e.g., on-line evaluation (stars))



Linear Classification Models

- Simple scenario: disjoint classes
- Divide input space into decision regions
- Linear model decision surface are linear function of input x
 - $D-1$ dimension hyperplane in D dimensional input space
- Linearly separable:
 - Classes in a dataset can be separated exactly using linear decision surface



Probabilistic Models

- Converting Regression to Class Output
- Target variables t now represent class labels
- Binary class representation is convenient
 - Two classes: $t \in \{0,1\}$, $t=1$ represent C_1 ; $t=0$ represent C_0
 - Interpret value of t as the probability that class is C_1
 - Probabilities need to take value between 0 and 1
 - For $K>2$, use 1-of- K coding scheme
 - E.g. $K=5$, $t = (0,1,0,0,0)^T$ represents class 2
 - Value t_k interpreted as probability of class C_k



Two Approaches to Classification

- Discriminant function
 - Directly assign x to a specific class
 - E.g. Fisher's linear discriminant function, perceptron
- Probabilistic models
 - Model $p(C_k|x)$ and make optimal decisions
- Probabilistic models separating inference from decision
 - Accommodate different loss function
 - Compensate for unbalanced data
 - Combine models



Probabilistic Classification Models

- Generative

- Model class conditional densities by $p(x|C_k)$ together with prior probabilities $p(C_k)$
- $p(C_k|x) \propto p(x|C_k)p(C_k)$

- Discriminative

- Directly model conditional probabilities $p(C_k|x)$



Converting Linear Regression Models to Linear Classification Models

- Linear regression model $y(x, w)$ is a linear function of parameters w
- In simplest case also a linear function of variable x
 - $y(x) = w^T x + w_0$
- For classification we wish to obtain a discrete output or posterior probabilities in range $(0, 1)$
- Use generalization of this model $y(x) = f(w^T x + w_0)$
 - $f(\cdot)$ is an activation function
 - $f^{-1}(\cdot)$ is a link function



Generalized Linear Model

- $y(x) = f(w^T x + w_0)$
- The decision surface correspond to $y(x) = \text{constant}$ or $w^T x + w_0 = \text{constant}$
- So decision boundaries are linear even if $f(\cdot)$ is nonlinear
 - → Generalized Linear Model
- No longer linear w.r.t to w due to $f(\cdot)$
 - Lead to more complex models for classification than regression



DISCRIMINANT FUNCTIONS


Linear Discriminant Function

- Assigns input vector x to one of K classes denoted by C_k
- Restrict attention to linear discriminants (decision hyperplanes)
- Two-class linear discriminant function:
 - $y(x) = w^T x + w_0$
 - w is a weight vector and w_0 is bias
 - Assign x to C_1 if $y(x) \geq 0$, else C_2




Geometry of Linear Discriminant Functions

- $y(x) = w^T x + w_0$
- Decision boundary: $y(x) = 0$
- For x_a and x_b on the boundary, $y(x_a) - y(x_b) = 0, \Rightarrow w^T(x_a - x_b) = 0$
- $\Rightarrow w$ is orthogonal to every vector on surface.
- Given an arbitrary vector x , let θ be the angle between x and w , then
$$\cos \theta = \frac{x \cdot w}{\|x\| \|w\|}$$
- The projection length of x on w is
$$\|x\| \cos \theta = \frac{w^T x}{\|w\|} = \frac{y(x) - w_0}{\|w\|} = \frac{y(x)}{\|w\|} + \frac{-w_0}{\|w\|}$$





Geometry of Linear Discriminant Functions

- $y(x) = w^T x + w_0$
- $y(x)$ gives signed measure of perpendicular distance r of point x to decision surface ($r = \frac{y(x)}{\|w\|}$)
- w_0 sets distance of origin to surface ($\frac{y(0)}{\|w\|} = \frac{w_0}{\|w\|}$)
- With dummy input $x_0 = 1$, and $\omega = (w_0, w^T)^T$, $z = (x_0, x^T)^T$, then $y(x) = \omega^T z$
 - Pass through origin in augmented $D+1$ dimensional space



Multiple Classes with 2-class classifiers

- One-versus-the-rest
- Build a K class discriminant using $K-1$ classifiers
- One-versus-one
- Build $K(K-1)/2$ binary discriminant classifiers, one for each pair



Both result in ambiguous regions of input space




Multiple Classes with K Discriminants

- Consider a single K class discriminants of the form
 - $y_k(x) = w_k^T x + w_{k0}$, $k = 1, 2, \dots, K$
 - Assign a point x to class C_k if $y_k(x) > y_j(x)$ for all $j \neq k$
- Decision boundary between class C_k and C_j is given by $y_k(x) = y_j(x)$
 - Corresponds to D-1 dimensional hyperplane defined by
$$(w_k - w_j)^T x + (w_{k0} - w_{j0}) = 0$$
 - Same form as the decision boundary for 2-class case
 - Decision regions of such a discriminant are always singly connected and convex
 - Proof on next slide



Convexity of Decision Regions

- Consider two points x_a and x_b both in decision region R_k
- Any point \hat{x} on line connecting x_a and x_b can be expressed as $\hat{x} = \lambda x_a + (1 - \lambda)x_b$, where $0 \leq \lambda \leq 1$.
- From linearity of discriminant functions $y_k(x) = w_k^T x + w_{k0}$
- Combining the two, we have
$$y_k(\hat{x}) = \lambda y_k(x_a) + (1 - \lambda)y_k(x_b)$$
- Because x_a and x_b lie inside R_k , it follows that $y_k(x_a) > y_j(x_a)$ and $y_k(x_b) > y_j(x_b)$ for all $j \neq k$.
- Hence \hat{x} also lies inside R_k
- Thus R_k is singly-connected and convex



Learning the Parameters of Linear Discriminant Functions

- Three methods
 - Least squares
 - Fisher's Linear Discriminant (omitted)
 - Perceptrons



Least Square for Classification

- Simple closed-form solution exists
- Each C_k , $k = 1, 2, \dots, K$, is described by its own linear model
$$y_k(x) = w_k^T x + w_{k0}$$
- Create augmented vector $x = (1, x^T)^T$ and
$$w_k = (w_{k0}, w_k^T)^T \text{ (length is } D+1 \text{)}$$
- Grouping into vector notation $y(x) = W^T x$
 - W is the parameter matrix whose k th column is a $D+1$ dimensional vector (including bias)
 - $W = [w_1 \ w_2 \ \dots \ w_K]$
 - $y(x)$ is K by 1
- New input vector x is assigned to class for which output $y_k = w_k^T x$ is largest
- Determine W by minimizing squared error



Learning Parameters Using Least Squares

- Training data $\{x_n, t_n\}$, $n = 1, 2, \dots, N$.
 t_n is a column vector of K dimensions using 1-of- K form

- Define matrices

$$T = \begin{bmatrix} t_1^T \\ t_2^T \\ \vdots \\ t_N^T \end{bmatrix}, X = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_N^T \end{bmatrix}, \text{Error vector} = \overset{N \times (D+1)}{X} \overset{(D+1) \times K}{W} - \overset{N \times K}{T}$$



- Let $B = XW - T = \begin{bmatrix} e_1^T \\ e_2^T \\ \vdots \\ e_N^T \end{bmatrix}$, we are trying to minimize the sum of squared errors $\sum_{j=1}^K \sum_{i=1}^N e_{ij}^2$.



Learning Regression Parameters

- We have k vectors to learn: $W = [w_1 \ w_2 \ \dots \ w_K]$
- Each w_i controls the i – th dimension of the output.
- All of the K vectors can be learned separately.
- That is, we can learn K separate regression models, each use one of the K columns in T as the outcome vector in the linear regression models we introduced before.
- To predict the output class for an input x , compute $y_i = w_i^T x$ for $i = 1, 2, \dots, K$, and
- classify x to class k for which y_k is the maximum.


[Limitation] Least Square is Sensitive to Outliers




- Sum of squared errors penalized predictions that are “too correct” (magenta lines)
 - Long way from decision boundary
 - Logistic regression (green lines) does not have this problem
 - SVMs have an hinge error function that have a similar shape to that of the logistic regression

Disadvantages of Least Squares

Ordinary
regression



Logistic
regression



- Regions assigned to green class is too small, mostly misclassified
- Logistic regression performs well
- Disadvantage of least square:
 - Lack robustness to outliers
 - Gaussian assumption may not be a good choice for classification problems
 - Binary target values have a distribution far from Gaussian



Perceptron Algorithm

- Two-class model
 - Input vector x transformed by a fixed nonlinear function to give feature vector $\phi(x)$
 - $y(x) = f(w^T \phi(x))$
where non-linear activation $f(\cdot)$ is a step function
 - $$f(a) = \begin{cases} +1 & \text{if } a \geq 0 \\ -1 & \text{if } a < 0 \end{cases}$$
- Use a target coding scheme
 - $t=1$ for class C_1 , and $t=-1$ for C_2 similar to the activation function



Perceptron Error Function

- Error function: **number of misclassifications**
- This error function is a piecewise constant function of w with discontinuities (unlike regression)
- Hence no closed form solution



Perceptron Criterion

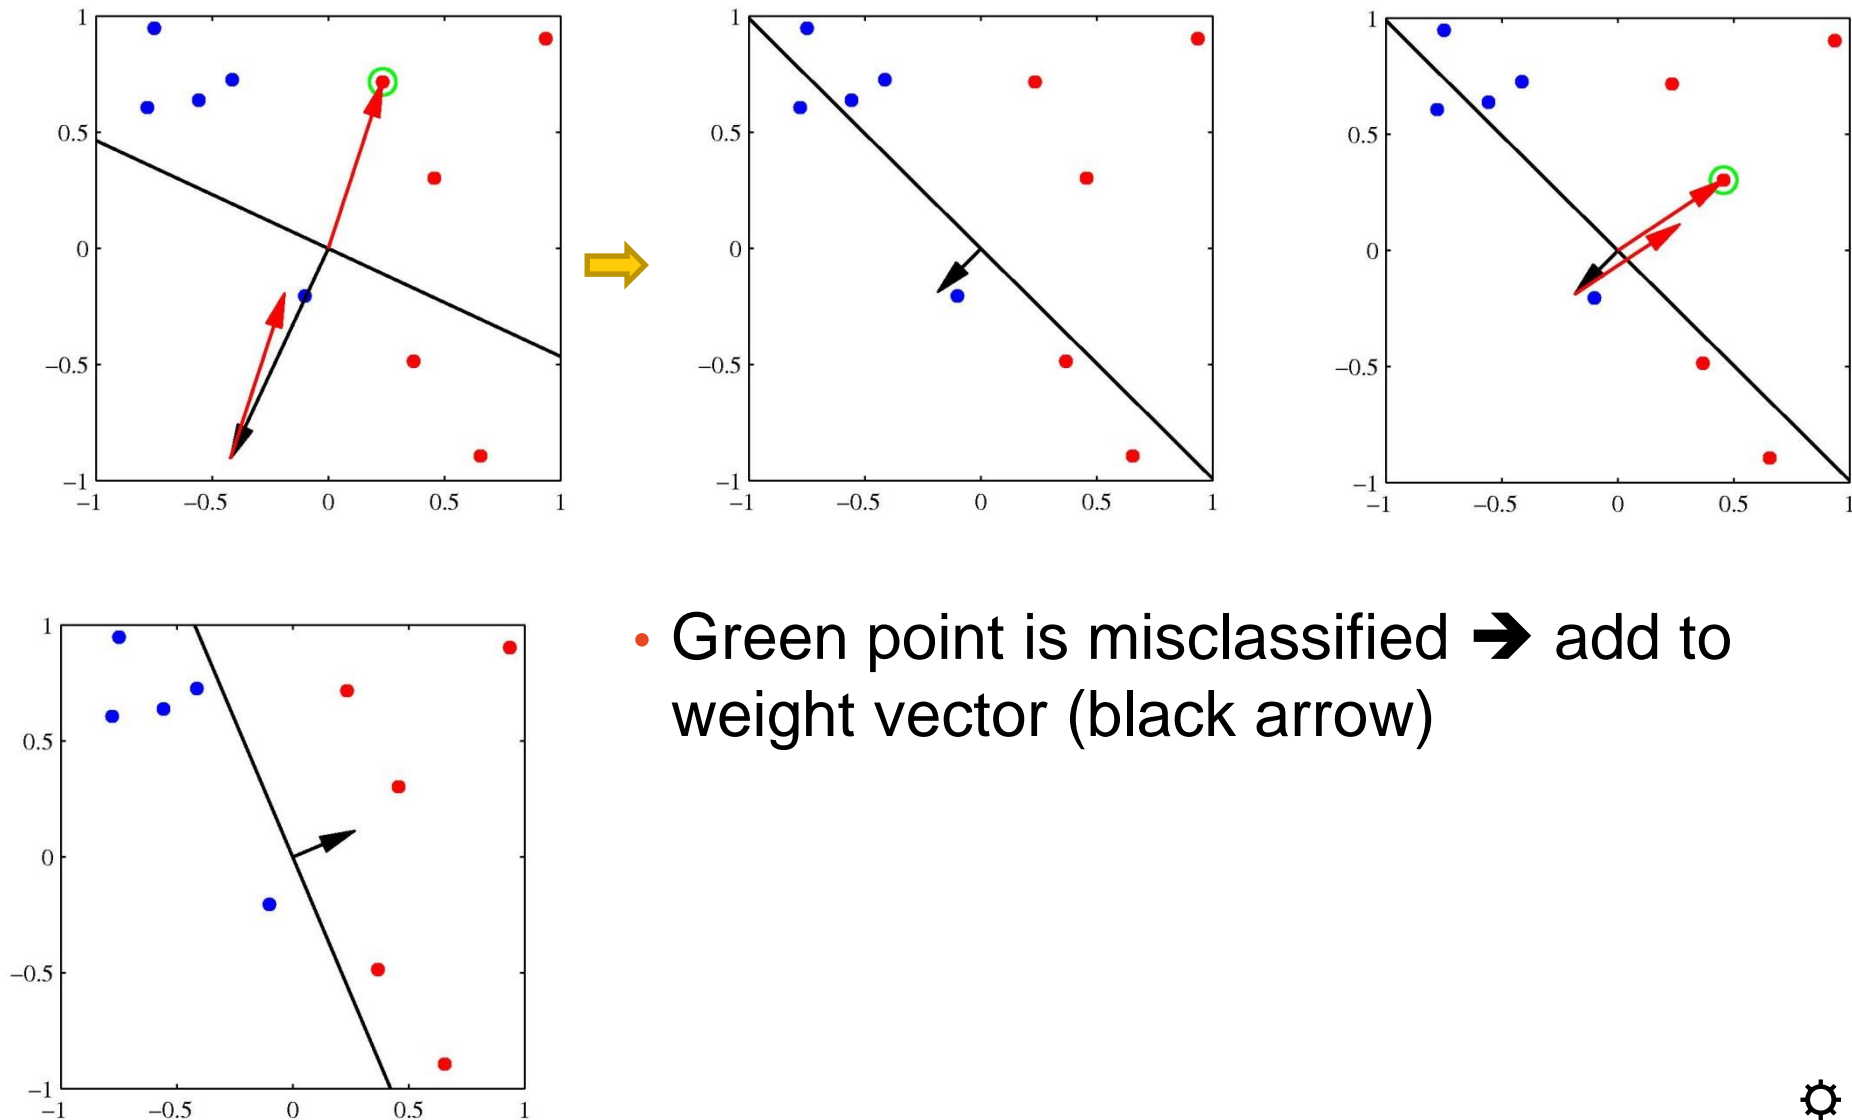
- Seek w such that $x_n \in C_1$ will have $w^T x_n \geq 0$ where as $x_n \in C_2$ will have $w^T x_n < 0$
- Target value $t \in \{-1, 1\}$. $C_1 \rightarrow t=1$, all patterns need to satisfy $w^T \phi(x_n)t_n > 0$
- For each misclassified sample, perceptron criterion tried to minimize
$$E_p(w) = - \sum_{n \in \mathbf{M}} w^T \phi_n t_n$$
- \mathbf{M} denotes set of all misclassified patterns and $\phi_n = \phi(x_n)$

Perceptron Algorithm

- Error function $E_p(w) = -\sum_{n \in M} w^T \phi_n t_n \equiv \sum_{n \in M} E_n$
- **Stochastic gradient descent**
 - a.k.a Sequential gradient descent
 - Loop through M , for each data point n , change in weight is given by
 - $w^{\tau+1} = w^{\tau} - \eta \nabla E_n(w) = w^{\tau} + \eta \phi_n t_n$
 - $\nabla E_n(w) = \frac{\partial E_n(w)}{\partial w} = -\phi_n t_n$
 η is learning rate (set to 1), τ is a step index
- The algorithm
 - Cycle through the training patterns in turn
 - If incorrectly classified for class C_1 , add to weight vector
 - If incorrectly classified for class C_2 , subtract from weight vector




Perceptron Learning



History of Perceptron



- Perceptron invented by Frank Rosenblatt
- Mark 1 perceptron hardware (patching board, allowing different configurations)



Disadvantages of Perceptrons

- Does not converge if classes not linearly separable
- Does not provide probabilistic output
- Not readily generalized to $K > 2$ classes



Comparison of Parameter Learning Approaches

- Least Squares
 - Not robust to outliers, model close to target values
- Fisher's linear discriminant (omitted)
 - Separation determined by within-class and between class variance
 - Can be seen as a special case of least square
- Perceptrons
 - Does not converge if classes not linearly separable
 - Does not provide probabilistic output
 - Not readily generalized to $K > 2$ classes



QUESTIONS?
