

LINEAR MODELS FOR CLASSIFICATION (PART 2)

Hsin-Min Lu

盧信銘

台大資管系

PROBABILISTIC DISCRIMINATIVE MODELS



Topics in Linear Classification using Probabilistic Discriminative Models


- Nonlinear basis functions in linear classification
- Logistic Regression
 - Two-class, Multi-class
 - Parameter estimation
 - Maximum Likelihood
 - Iterative Reweighted Least Squares
- Probit Regression



Nonlinear Basis Functions in Linear Models

- Although we use linear classification models
 - Linear-separability **in feature space** does not imply linear-separability in **input space**
 - Developing useful features to enable linear separable feature space is called “**feature engineering**.”

Original Input Space




not linearly separable

**Nonlinear
transformation
of inputs using vector of
basis functions $\phi(x)$**



Feature Space (ϕ_1, ϕ_2)



linearly separable



2. Logistic Regression


- Feature vector ϕ , two-classes C_1 and C_2
- A *posteriori* probability $p(C_1|\phi)$ can be written as

$$p(C_1|\phi) = y(\phi) = \sigma(w^T \phi)$$


where ϕ is a M -dimensional feature vector
 $\sigma(\cdot)$ is the logistic sigmoid function

- Goal is to determine the M parameters
- Known as logistic regression in statistics
 - Although a model for classification rather than for regression

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$



Logistic Sigmoid



- $\sigma(a) = \frac{1}{1+\exp(-a)}$
- $1 - \sigma(a) = \frac{1+\exp(-a)-1}{1+\exp(-a)} = \frac{\exp(-a)}{1+\exp(-a)} = \frac{1}{\exp(a)+1} = \sigma(-a)$
- Symmetric: $\sigma(-a) = 1 - \sigma(a)$
- $a = \ln\left(\frac{\sigma}{1-\sigma}\right)$ is the log odds ratio
- $\ln\left(\frac{\sigma}{1-\sigma}\right) = \ln \frac{p(C_1|x)}{p(C_2|x)}$
- $\frac{\partial \sigma(a)}{\partial a} = (-1)(1 + \exp(-a))^{-2}(-\exp(-a))$
- $= \frac{\exp(-a)}{[1+\exp(-a)]^2} = \sigma(a) \frac{\exp(-a)}{1+\exp(-a)} = \sigma(a)(1 - \sigma(a))$

Fewer Parameters in Linear Discriminative Model

- Discriminative approach (Logistic Regression)
 - For M -dim space: M adjustable parameters
- Generative based on Gaussians (Bayes/NB) (Omitted)
 - $2M$ parameters for mean
 - $M(M+1)/2$ parameters for shared covariance matrix
 - Two class priors
 - Total of $M(M+5)/2 + 1$ parameters
 - Grows quadratically with M
 - If features assumed independent (naïve Bayes) still needs $M+3$ parameters



Determining Logistic Regression parameters

- Maximum Likelihood Approach for Two classes

Data set $\{\phi_n, t_n\}$

where $t_n \in \{0,1\}$ and $\phi_n = \phi(x_n), n = 1, \dots, N$

Since t_n is binary we can use Bernoulli

Let y_n be the probability that $t_n = 1$

- Likelihood function associated with N observations

$$p(t|w) = \prod_{n=1}^N y_n^{t_n} \{1 - y_n\}^{1-t_n}$$

- where $t = (t_1, \dots, t_N)^T$, and
- $y_n = p(C_1|\phi_n)$



Error Function for Logistic Regression

Likelihood function

$$p(t|w) = \prod_{n=1}^N y_n^{t_n} \{1 - y_n\}^{1-t_n}$$

Error function is the negative of the log-likelihood

$$\begin{aligned} E(w) &= -\ln p(t|w) \\ &= -\sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln(1 - y_n)\} \end{aligned}$$



Gradient of Error Function

- Error function:
- $E(w) = -\ln p(t|w) = -\sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln(1 - y_n)\}$
- where $y_n = \sigma(w^T \phi_n)$
- Want to compute $\nabla E(w) = \frac{\partial E(w)}{\partial w}$
- Recall: $\frac{\partial \sigma(a)}{\partial a} = \sigma(a)(1 - \sigma(a))$,
- Let $z = z_1 + z_2$,
- where $z_1 = t \ln \sigma(w^T \phi)$, $z_2 = (1 - t) \ln[1 - \sigma(w^T \phi)]$.
- $\frac{\partial z_1}{\partial w} = \frac{t \sigma(w^T \phi)(1 - \sigma(w^T \phi))\phi}{\sigma(w^T \phi)} = t[1 - \sigma(w^T \phi)]\phi$
- $\frac{\partial z_2}{\partial w} = \frac{(1-t)(-1)\sigma(w^T \phi)(1 - \sigma(w^T \phi))\phi}{1 - \sigma(w^T \phi)} = (t - 1)\sigma(w^T \phi)\phi$
- $\frac{\partial z}{\partial w} = \frac{\partial z_1}{\partial w} + \frac{\partial z_2}{\partial w} = (t - \sigma(w^T \phi))\phi$



Gradient of Error Function (Cont'd.)


- Error function:
- $E(w) = -\ln p(t|w) = -\sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln(1 - y_n)\}$
- where $y_n = \sigma(w^T \phi_n)$
- Let $z = z_1 + z_2$,
- where $z_1 = t \ln \sigma(w^T \phi)$, $z_2 = (1 - t) \ln[1 - \sigma(w^T \phi)]$.
- $\frac{\partial z}{\partial w} = \frac{\partial z_1}{\partial w} + \frac{\partial z_2}{\partial w} = (t - \sigma(w^T \phi))\phi$
- $\nabla E(w) = \frac{\partial E(w)}{\partial w} = \sum_{n=1}^N (y_n - t_n)\phi_n$
- Contribution to gradient by data point n is error between target t_n and prediction $y_n = \sigma(w^T \phi_n)$ times basis ϕ_n



Finding the Minimal Value of $E(w)$

- It is clear from the previous slide that we cannot not solve $\nabla E(w) = \sum_{n=1}^N (y_n - t_n) \phi_n = 0$ for w directly.
- Need numerical algorithm for this task.
- Solving for $\nabla E(w) = 0$ is sometimes called the **root finding** problem.
- One well known approach is called the *Newton-Raphson* Method.
- Start with x_n , use the tangent line at $f(x_n)$ to determine x_{n+1} , and repeat.

Newton's Method




$$f'(x_n) = \frac{\text{rise}}{\text{run}} = \frac{\Delta y}{\Delta x} = \frac{f(x_n) - 0}{x_n - x_{n+1}}.$$

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}.$$

Since we are solving for the derivative of $E(w)$
Hessian of $E(w)$ is needed

Newton-Raphson Method (One Dimensional)

- Example: Given $f(x) = x^5 + x - 1$
- Use $x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$ to approximate the root
- $f'(x) = 5x^4 + 1$, we start from $x_1 = 1$,
- starting from $x = 1$
- step 1 : 0.5
- step 2 : -0.6666667
- step 3 : -0.7681159
- step 4 : -0.7551625
- step 5 : -0.7548778
- step 6 : -0.7548777
- step 7 : -0.7548777
- step 8 : -0.7548777
- step 9 : -0.7548777
- step 10 : -0.7548777



Iterative Reweighted Least Squares (IRLS)

- Efficient approximation using *Newton-Raphson* iterative optimization
- $w^{(new)} = w^{(old)} - H^{-1} \nabla E(w)$
 - where H is the Hessian matrix
- $H = \frac{\partial}{\partial w^T} \left(\frac{\partial E(w)}{\partial w} \right) = \frac{\partial E(w)}{\partial w \partial w^T} = \nabla \nabla E(w)$
- Elements in H are the second derivatives of $E(w)$ with respect to the components of w



Two applications of IRLS

- IRLS is applicable to both Linear Regression and Logistic Regression
- We discuss both, for each we need
 1. Error function $E(w)$
 - Linear Regression: Sum of Squared Errors
 - Logistic Regression: Bernoulli Likelihood Function
 2. Gradient $\nabla E(w)$
 3. Hessian $H = \nabla \nabla E(w)$
 4. Newton-Raphson update
$$w^{(new)} = w^{(old)} - H^{-1} \nabla E(w)$$



IRLS for Linear Regression

- Recall the linear regression model is:

- $y(x, w) = \sum_{j=0}^{M-1} w_j \phi_j(x) = w^T \phi(x)$

- Data set: $X = \{x_n, t_n\} \ n = 1, \dots, N$

- Error Function: Sum of Squared Errors

$$E(w) = \frac{1}{2} \sum_{n=1}^N \{t_n - w^T \phi(x_n)\}^2 = \frac{1}{2} (\mathbf{t} - \Phi \mathbf{w})^T (\mathbf{t} - \Phi \mathbf{w})$$

$$= \frac{1}{2} [\mathbf{t}^T \mathbf{t} - 2 \mathbf{w}^T \Phi^T \mathbf{t} + \mathbf{w}^T \Phi^T \Phi \mathbf{w}]$$

$$\Phi = \begin{pmatrix} \phi_0(x_1) & \phi_1(x_1) & \dots & \phi_{M-1}(x_1) \\ \phi_0(x_2) & & & \\ & & & \\ \phi_0(x_N) & & & \phi_{M-1}(x_N) \end{pmatrix}$$

- Gradient of Error Function is:

$$\nabla E(w) = \Phi^T \Phi \mathbf{w} - \Phi^T \mathbf{t}$$

Φ is the $N \times M$ design matrix
whose n^{th} row is given by ϕ_n^T

- Hessian is:

$$H = \nabla \nabla E(w) = \sum_{n=1}^N \phi_n \phi_n^T = \Phi^T \Phi$$



4. Newton-Raphson for Linear Regression

$$w^{(new)} = w^{(old)} - H^{-1} \nabla E(w)$$

Substituting: $H = \Phi^T \Phi$ and $\nabla E(w) = \Phi^T \Phi w - \Phi^T \mathbf{t}$

$$\begin{aligned} w^{(new)} &= w^{(old)} - (\Phi^T \Phi)^{-1} \{ \Phi^T \Phi w^{(old)} - \Phi^T \mathbf{t} \} \\ &= (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \end{aligned}$$

which is the standard least squares solution

Since it is independent of w , Newton-Raphson gives exact solution in one step



IRLS for Logistic Regression

- Posterior probability of class C_1 is

$$p(C_1|\phi) = y(\phi) = \sigma(w^T \phi)$$

- Likelihood Function

- For data set $\{\phi_n, t_n\}$, $t_n \in \{0,1\}$, $\phi_n = \phi(x_n)$

- $p(t|w) = \prod_{n=1}^N y_n^{t_n} \{1 - y_n\}^{1-t_n}$

- Error Function

$$E(w) = - \sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln(1 - y_n)\}$$

$$y_n = \sigma(w^T \phi_n)$$



IRLS for Logistic Regression

- Gradient of Error Function:

$$\nabla E(w) = \sum_{n=1}^N (y_n - t_n) \phi_n = \Phi^T (y - t)$$

- $y_n = \sigma(w^T \phi_n)$
- Hessian: $H = \nabla \nabla E(w) = \sum_{n=1}^N y_n (1 - y_n) \phi_n \phi_n^T = \Phi^T R \Phi$
- R is NxN diagonal matrix with elements: $R_{nn} = y_n (1 - y_n)$
- Hessian is not constant and depends on w through R
- Since H is positive-definite (i.e., for arbitrary u, $u^T H u > 0$), error function is a concave function of w and so has a unique minimum



IRLS for Logistic Regression

- Newton-Raphson update:

$$w^{(new)} = w^{(old)} - H^{-1} \nabla E(w)$$

- Substituting $H = \Phi^T R \Phi$ and $\nabla E(w) = \Phi^T (y - t)$
 - $w^{(new)} = w^{(old)} - (\Phi^T R \Phi)^{-1} \Phi^T (y - t)$
 - $= (\Phi^T R \Phi)^{-1} \{ \Phi^T R \Phi w^{(old)} - \Phi^T (y - t) \}$
 - $= (\Phi^T R \Phi)^{-1} \Phi^T R \{ \Phi w^{(old)} - R^{-1} (y - t) \}$
 - $= (\Phi^T R \Phi)^{-1} \Phi^T R z$
- z is a N -dimensional vector with elements
- $z = \Phi w^{(old)} - R^{-1} (y - t)$

Update formula is a set of normal equations

Since Hessian depends on w

Apply them iteratively each time using the new weight vector



Question

- What may go wrong if we have a lot of features?

Multi-class Logistic Regression

- Work with soft-max function instead of logistic sigmoid
- $P(C_k|\phi) = y_k(\phi) = \frac{\exp(a_k)}{\sum_j \exp(a_j)}$
- where $a_k = w_k^T \phi$



Multi-class Likelihood Function

- 1-of K Coding Scheme
 - For feature vector ϕ_n , target vector t_n belonging to class C_k is a binary vector with all elements zero except for element k

$$p(T|w_1, \dots, w_K) = \prod_{n=1}^N \prod_{k=1}^K p(C_k|\phi_n)^{t_{nk}} = \prod_{n=1}^N \prod_{k=1}^K y_{nk}^{t_{nk}}$$

- where $y_{nk} = y_k(\phi_n)$
- subject to (s.t.) $\sum_k t_{nk} = 1$ (because of 1-of-K coding)
- T is a $N \times K$ matrix of elements with elements t_{nk}



Multi-class Error Function

- Error Function: negative log-likelihood

$$E(w_1, \dots, w_k) = -\ln p(T|w_1, \dots, w_k) = -\sum_{n=1}^N \sum_{k=1}^K t_{nk} \ln y_{nk}$$

- Derivatives of soft-max:

- $y_k(\phi) = \frac{\exp(a_k)}{\sum_j \exp(a_j)}$, $a_k = w_k^T \phi$

- $\frac{\partial y_k}{\partial a_j} = y_k(I_{kj} - y_j)$, where I_{kj} are elements of the identity matrix. [Exercise]



Multi-class Error Function

- Error Function: negative log-likelihood

$$E(w_1, \dots, w_k) = -\ln p(T|w_1, \dots, w_k) = -\sum_{n=1}^N \sum_{k=1}^K t_{nk} \ln y_{nk}$$

- $y_{nk} = \frac{\exp(a_{nk})}{\sum_j \exp(a_{nj})}$, $a_{nk} = w_k^T \phi_n$

- Want: $\frac{\partial E(\cdot)}{\partial w_j} = \sum_{n=1}^N \sum_{k=1}^K \frac{-\partial t_{nk} \ln y_{nk}}{\partial y_{nk}} \frac{\partial y_{nk}}{\partial a_{nj}} \frac{\partial a_{nj}}{\partial w_j} =$

- $= \sum_{n=1}^N \sum_{k=1}^K -\frac{t_{nk}}{y_{nk}} [\textcolor{red}{y_{nk}} (I_{jk} - \textcolor{red}{y_{nj}})] \phi_n$

- $= \sum_{n=1}^N \sum_{k=1}^K (y_{nj} t_{nk} - I_{jk} t_{nk}) \phi_n$

- $= \sum_{n=1}^N \underbrace{(y_{nj} - t_{nj})}_{\text{Prediction error of the j-th outcome}} \phi_n \quad (\text{since } \sum_k t_{nk} = 1)$

Prediction error of the j-th outcome



IRLS Algorithm for Multi-class

- Hessian matrix comprises blocks of size $M \times M$
 - Block j,k is given by

$$\nabla_{w_k} \nabla_{w_j} E(w_1, \dots, w_k) = - \sum_{n=1}^N y_{nk} (I_{kj} - y_{nj}) \phi_n \phi_n^T$$

- Hessian matrix is positive-definite, therefore error function has a unique minimum
 - Note the total size of H is $MK \times MK$
- Batch Algorithm based on Newton-Raphson



3. Probit Regression

- Logistic transformation is good for exponential family
- Not suitable for some types of probability-based inference models (e.g. Gaussian Mixture)
- Alternative discriminative model is based on probit function (which is the CDF of a zero-mean Gaussian)
 - Note that a CDF also goes between 0 and 1



Probit Activation Function

- Two-class case, Generalized Linear Model

$$p(t = 1|a) = f(a)$$

- where $a = w^T \phi$ and $f(\cdot)$ is the activation function

- Consider stochastic (noisy) threshold model

- For input ϕ_n , evaluate $a_n = w^T \phi_n$ and draw a random variable $\theta \sim N(\theta; 0, 1)$.

- Assign target value as
$$\begin{cases} t_n = 1 & \text{if } a_n \geq \theta \\ t_n = 0 & \text{otherwise} \end{cases}$$


- Then
$$p(t_n = 1|a_n) = p(\theta \leq a_n) = \int_{-\infty}^{a_n} N(\theta; 0, 1) d\theta = \Phi(a_n)$$



Probit Function

It is the CDF of a zero-mean unit-variance Gaussian

$$\Phi(a) = \int_{-\infty}^a N(\theta|0,1)d\theta$$



It has a *sigmoidal* shape & related to the *erf* function which is usually tabulated

$$\text{erf}(a) = \frac{2}{\sqrt{\pi}} \int_0^a \exp\left(-\frac{\theta^2}{2}\right) d\theta$$

← It represents the probability that the error lies between $\pm a$

With the relationship

$$\Phi(a) = \frac{1}{2} \left\{ 1 + \frac{1}{\sqrt{2}} \text{erf}(a) \right\}$$



Probit Regression

- Generalized Linear Model based on Probit activation function
- Parameters determined using maximum likelihood
- Results similar to Logistic Regression
 - Probit is more sensitive to outliers

